

Dueling Network Architectures for Deep Reinforcement Learning (ICML 2016)

Yoonho Lee

Department of Computer Science and Engineering
Pohang University of Science and Technology

October 11, 2016

Outline

Reinforcement Learning

- Definition of RL

- Mathematical formulations

Algorithms

- Neural Fitted Q Iteration

- Deep Q Network

- Double Deep Q Network

- Prioritized Experience Replay

- Dueling Network

Outline

Reinforcement Learning

- Definition of RL

- Mathematical formulations

Algorithms

- Neural Fitted Q Iteration

- Deep Q Network


- Double Deep Q Network

- Prioritized Experience Replay

- Dueling Network


Definition of RL

general setting of RL



Definition of RL

atari setting



Outline

Reinforcement Learning

Definition of RL

Mathematical formulations

Algorithms

Neural Fitted Q Iteration

Deep Q Network

Double Deep Q Network

Prioritized Experience Replay

Dueling Network

Mathematical formulations

objective of RL

Definition

Return G_t is the cumulative discounted reward from time t

$$G_t = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots$$

Mathematical formulations

objective of RL

Definition

Return G_t is the cumulative discounted reward from time t

$$G_t = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots$$

Definition

A policy π is a function that selects actions given states

$$\pi(s) = a$$

- The goal of RL is to find π that maximizes G_0

Mathematical formulations

Q-Value

$$G_t = \sum_{i=0}^{\infty} \gamma^i r_{t+i+1}$$

Definition

The action-value (Q-value) function $Q_{\pi}(s, a)$ is the expectation of G_t under taking action a , and then following policy π afterwards

$$Q_{\pi}(s, a) = \mathbb{E}_{\pi}[G_t | S_t = s, A_t = a, A_{t+i} = \pi(S_{t+i}) \forall i \in \mathbb{N}]$$

- "How good is action a in state s ?"

Mathematical formulations

Optimal Q-value

Definition

The optimal Q-value function $Q_*(s, a)$ is the maximum Q-value over all policies

$$Q_*(s, a) = \max_{\pi} Q_{\pi}(s, a)$$

Mathematical formulations

Optimal Q-value

Definition

The optimal Q-value function $Q_*(s, a)$ is the maximum Q-value over all policies

$$Q_*(s, a) = \max_{\pi} Q_{\pi}(s, a)$$

Theorem

There exists a policy π_* such that $Q_{\pi_*}(s, a) = Q_*(s, a)$ for all s, a

- Thus, it suffices to find Q_*

Mathematical formulations

Q-Learning

Bellman Optimality Equation

$Q_*(s, a)$ satisfies the following equation:

$$Q_*(s, a) = R(s, a) + \gamma \sum_{s'} P(s'|s, a) \max_{a'} Q_*(s', a')$$

Mathematical formulations

Q-Learning

Bellman Optimality Equation

$Q_*(s, a)$ satisfies the following equation:

$$Q_*(s, a) = R(s, a) + \gamma \sum_{s'} P(s'|s, a) \max_{a'} Q_*(s', a')$$

Q-Learning

Let a be ϵ -greedy w.r.t. Q , and a' be optimal w.r.t. Q . Q converges to Q_* if we iteratively apply the following update:

$$Q(s, a) \leftarrow \alpha(R(s, a) + \gamma Q(s', a')) + (1 - \alpha)Q(s, a)$$

Mathematical formulations

Other approaches to RL

Value-Based RL

- ▶ Estimate $Q_*(s, a)$
- ▶ Deep Q Network

Policy-Based RL

- ▶ Search directly for optimal policy π_*
- ▶ DDPG, TRPO...

Model-Based RL

- ▶ Use the (learned or given) transition model of environment
- ▶ Tree Search, DYNA ...

Outline

Reinforcement Learning

- Definition of RL

- Mathematical formulations

Algorithms

- Neural Fitted Q Iteration

- Deep Q Network

- Double Deep Q Network

- Prioritized Experience Replay

- Dueling Network

Neural Fitted Q Iteration

```
NFQ_main() {  
  input: a set of transition samples  $D$ ; output: Q-value function  $Q_N$   
  k=0  
  init_MLP()  $\rightarrow Q_0$ ;  
  Do {  
    generate_pattern_set  $P = \{(input^l, target^l), l = 1, \dots, \#D\}$  where:  
       $input^l = s^l, u^l$ ,  
       $target^l = c(s^l, u^l, s'^l) + \gamma \min_b Q_k(s'^l, b)$   
    Rprop_training( $P$ )  $\rightarrow Q_{k+1}$   
    k:= k+1  
  } WHILE ( $k < N$ )
```

Fig. 1. Main loop of NFQ

- Supervised learning on (s, a, r, s')

Neural Fitted Q Iteration

input

```
NFQ_main() {  
  input: a set of transition samples  $D$ ; output: Q-value function  $Q_N$   
  k=0  
  init_MLP()  $\rightarrow Q_0$ ;  
  Do {  
    generate_pattern_set  $P = \{(input^l, target^l), l = 1, \dots, \#D\}$  where:  
       $input^l = s^l, u^l$ ,  
       $target^l = c(s^l, u^l, s'^l) + \gamma \min_b Q_k(s'^l, b)$   
    Rprop_training( $P$ )  $\rightarrow Q_{k+1}$   
    k:= k+1  
  } WHILE ( $k < N$ )
```

Fig. 1. Main loop of NFQ

Neural Fitted Q Iteration

target equation

```
NFQ_main() {  
  input: a set of transition samples  $D$ ; output: Q-value function  $Q_N$   
  k=0  
  init_MLP()  $\rightarrow Q_0$ ;  
  Do {  
    generate_pattern_set  $P = \{(input^l, target^l), l = 1, \dots, \#D\}$  where:  
     $input^l = s^l, u^l,$   
     $target^l = c(s^l, u^l, s'^l) + \gamma \min_b Q_k(s'^l, b)$   
    Rprop_training( $P$ )  $\rightarrow Q_{k+1}$   
    k:= k+1  
  } WHILE ( $k < N$ )
```

Fig. 1. Main loop of NFQ

Neural Fitted Q Iteration

Shortcomings

- ▶ Exploration is independent of experience
- ▶ Exploitation does not occur at all
- ▶ Policy evaluation does not occur at all
- ▶ Even in exact(table lookup) case, not guaranteed to converge to Q_*

Outline

Reinforcement Learning

- Definition of RL

- Mathematical formulations

Algorithms

- Neural Fitted Q Iteration

- Deep Q Network

- Double Deep Q Network

- Prioritized Experience Replay

- Dueling Network

Deep Q Network

action policy

Algorithm 1: deep Q-learning with experience replay.

Initialize replay memory D to capacity N

Initialize action-value function Q with random weights θ

Initialize target action-value function \hat{Q} with weights $\theta^- = \theta$

For episode = 1, M **do**

 Initialize sequence $s_1 = \{x_1\}$ and preprocessed sequence $\phi_1 = \phi(s_1)$

For $t = 1, T$ **do**

 With probability ε select a random action a_t

 otherwise select $a_t = \operatorname{argmax}_a Q(\phi(s_t), a; \theta)$

 Execute action a_t in emulator and observe reward r_t and image x_{t+1}

 Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$

 Store transition $(\phi_t, a_t, r_t, \phi_{t+1})$ in D

 Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from D

 Set $y_j = \begin{cases} r_j & \text{if episode terminates at step } j+1 \\ r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$

 Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ with respect to the network parameters θ

 Every C steps reset $\hat{Q} = Q$

End For

End For

- Choose an ϵ -greedy policy w.r.t. Q

Deep Q Network

network freezing

Algorithm 1: deep Q-learning with experience replay.

Initialize replay memory D to capacity N

Initialize action-value function Q with random weights θ

Initialize target action-value function \hat{Q} with weights $\theta^- = \theta$

For episode = 1, M **do**

Initialize sequence $s_1 = \{x_1\}$ and preprocessed sequence $\phi_1 = \phi(s_1)$

For $t = 1, T$ **do**

With probability ε select a random action a_t

otherwise select $a_t = \operatorname{argmax}_a Q(\phi(s_t), a; \theta)$

Execute action a_t in emulator and observe reward r_t and image x_{t+1}

Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$

Store transition $(\phi_t, a_t, r_t, \phi_{t+1})$ in D

Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from D

Set $y_j = \begin{cases} r_j & \text{if episode terminates at step } j+1 \\ r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$

Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ with respect to the network parameters θ

Every C steps reset $\hat{Q} = Q$

End For

End For

- Get a copy of the network every C steps for stability

Deep Q Network

Algorithm 1: deep Q-learning with experience replay.

Initialize replay memory D to capacity N

Initialize action-value function Q with random weights θ

Initialize target action-value function \hat{Q} with weights $\theta^- = \theta$

For episode = 1, M **do**

 Initialize sequence $s_1 = \{x_1\}$ and preprocessed sequence $\phi_1 = \phi(s_1)$

For $t = 1, T$ **do**

 With probability ε select a random action a_t

 otherwise select $a_t = \operatorname{argmax}_a Q(\phi(s_t), a; \theta)$

 Execute action a_t in emulator and observe reward r_t and image x_{t+1}

 Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$

 Store transition $(\phi_t, a_t, r_t, \phi_{t+1})$ in D

 Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from D

 Set $y_j = \begin{cases} r_j & \text{if episode terminates at step } j+1 \\ r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$

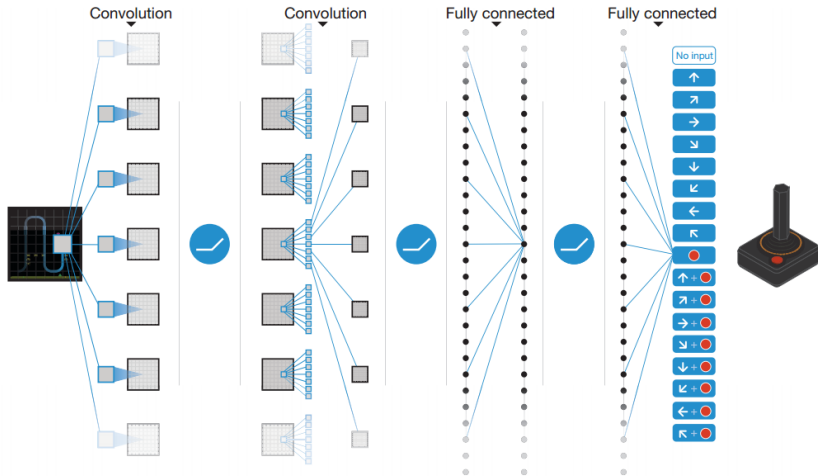
 Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ with respect to the network parameters θ

 Every C steps reset $\hat{Q} = Q$

End For


End For

Deep Q Network



Deep Q Network

performance



Deep Q Network

overoptimism

Algorithm 1: deep Q-learning with experience replay.

Initialize replay memory D to capacity N

Initialize action-value function Q with random weights θ

Initialize target action-value function \hat{Q} with weights $\theta^- = \theta$

For episode = 1, M **do**

 Initialize sequence $s_1 = \{x_1\}$ and preprocessed sequence $\phi_1 = \phi(s_1)$

For $t = 1, T$ **do**

 With probability ε select a random action a_t

 otherwise select $a_t = \operatorname{argmax}_a Q(\phi(s_t), a; \theta)$

 Execute action a_t in emulator and observe reward r_t and image x_{t+1}

 Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$

 Store transition $(\phi_t, a_t, r_t, \phi_{t+1})$ in D

 Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from D

 Set $y_j = \begin{cases} r_j & \text{if episode terminates at step } j+1 \\ r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$

 Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ with respect to the network parameters θ

 Every C steps reset $\hat{Q} = Q$

End For

End For

- What happens when we overestimate Q ?

Deep Q Network

Problems

- ▶ Overestimation of the Q function at any s spills over to actions that lead to $s \rightarrow$ Double DQN
- ▶ Sampling transitions uniformly from D is inefficient \rightarrow Prioritized Experience Replay

Outline

Reinforcement Learning

- Definition of RL

- Mathematical formulations

Algorithms

- Neural Fitted Q Iteration

- Deep Q Network

- Double Deep Q Network**

- Prioritized Experience Replay

- Dueling Network

Double Deep Q Network

target

We can write the DQN target as:

$$Y_t^{DQN} = R_{t+1} + \gamma Q(S_{t+1}, \underset{a}{\operatorname{argmax}} Q(S_{t+1}, a; \theta_t^-); \theta_t^-)$$

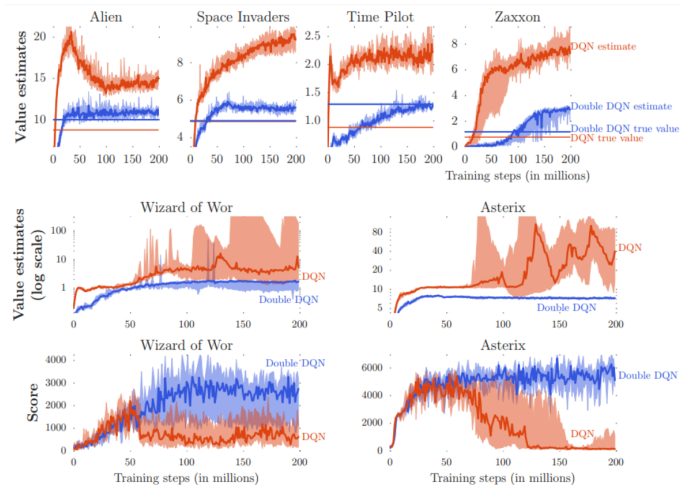
Double DQN's target is:

$$Y_t^{DDQN} = R_{t+1} + \gamma Q(S_{t+1}, \underset{a}{\operatorname{argmax}} Q(S_{t+1}, a; \theta_t^-); \theta_t^-)$$

This has the effect of decoupling action selection and action evaluation

Double Deep Q Network

performance



- Much stabler with very little change in code

Outline

Reinforcement Learning

- Definition of RL

- Mathematical formulations

Algorithms

- Neural Fitted Q Iteration

- Deep Q Network

- Double Deep Q Network

- Prioritized Experience Replay

- Dueling Network

Prioritized Experience Replay

Algorithm 1 Double DQN with proportional prioritization

```
1: Input: minibatch  $k$ , step-size  $\eta$ , replay period  $K$  and size  $N$ , exponents  $\alpha$  and  $\beta$ , budget  $T$ .
2: Initialize replay memory  $\mathcal{H} = \emptyset$ ,  $\Delta = 0$ ,  $p_1 = 1$ 
3: Observe  $S_0$  and choose  $A_0 \sim \pi_\theta(S_0)$ 
4: for  $t = 1$  to  $T$  do
5:   Observe  $S_t, R_t, \gamma_t$ 
6:   Store transition  $(S_{t-1}, A_{t-1}, R_t, \gamma_t, S_t)$  in  $\mathcal{H}$  with maximal priority  $p_t = \max_{i < t} p_i$ 
7:   if  $t \equiv 0 \pmod K$  then
8:     for  $j = 1$  to  $k$  do
9:       Sample transition  $j \sim P(j) = p_j^\alpha / \sum_i p_i^\alpha$ 
10:      Compute importance-sampling weight  $w_j = (N \cdot P(j))^{-\beta} / \max_i w_i$ 
11:      Compute TD-error  $\delta_j = R_j + \gamma_j Q_{\text{target}}(S_j, \arg \max_a Q(S_j, a)) - Q(S_{j-1}, A_{j-1})$ 
12:      Update transition priority  $p_j \leftarrow |\delta_j|$ 
13:      Accumulate weight-change  $\Delta \leftarrow \Delta + w_j \cdot \delta_j \cdot \nabla_\theta Q(S_{j-1}, A_{j-1})$ 
14:     end for
15:     Update weights  $\theta \leftarrow \theta + \eta \cdot \Delta$ , reset  $\Delta = 0$ 
16:     From time to time copy weights into target network  $\theta_{\text{target}} \leftarrow \theta$ 
17:   end if
18:   Choose action  $A_t \sim \pi_\theta(S_t)$ 
19: end for
```

- Update 'more surprising' experiences more frequently

Outline

Reinforcement Learning

- Definition of RL

- Mathematical formulations

Algorithms

- Neural Fitted Q Iteration


- Deep Q Network

- Double Deep Q Network

- Prioritized Experience Replay

- Dueling Network

Dueling Network



- The scalar approximates V , and the vector approximates A

Dueling Network

forward pass equation


The exact forward pass equation is:

$$Q(s, a; \theta, \alpha, \beta) = V(s; \theta, \beta) + (A(s, a; \theta, \alpha) - \max_{a'} A(s, a'; \theta, \alpha))$$

The following module was found to be more stable without losing much performance:

$$Q(s, a; \theta, \alpha, \beta) = V(s; \theta, \beta) + (A(s, a; \theta, \alpha) - \frac{1}{|A|} \sum_{a'} A(s, a'; \theta, \alpha))$$

Dueling Network



Dueling Network

performance

| | 30 no-ops | | Human Starts | |
|------------------|---------------|---------------|---------------|---------------|
| | Mean | Median | Mean | Median |
| Prior. Duel Clip | 591.9% | 172.1% | 567.0% | 115.3% |
| Prior. Single | 434.6% | 123.7% | 386.7% | 112.9% |
| Duel Clip | 373.1% | 151.5% | 343.8% | 117.1% |
| Single Clip | 341.2% | 132.6% | 302.8% | 114.1% |
| Single | 307.3% | 117.8% | 332.9% | 110.9% |
| Nature DQN | 227.9% | 79.1% | 219.6% | 68.5% |

- Achieves state of the art in the Atari domain among DQN algorithms

Dueling Network

Summary

- ▶ Since this is an improvement only in network architecture, methods that improve DQN(e.g. Double DQN) are all applicable here as well
- ▶ Solves problem of V and A typically being of different scale
- ▶ Updates Q values more frequently than a single-stream DQN, where only a single Q value is updated for each observation
- ▶ Implicitly splits the credit assignment problem into a recursive binary problem of “now or later”

Dueling Network

Shortcomings

- ▶ Only works for $|A| < \infty$
- ▶ Still not able to solve tasks involving long-term planning
- ▶ Better than DQN, but sample complexity is still high
- ▶ ϵ -greedy exploration is essentially random guessing